

MPLS to SASE Migration & Co-existence

Nobody migrates a WAN in one weekend. This use case shows how Cato lets migrated and not-yet-migrated sites keep talking to each other throughout a site-by-site MPLS exit — the Cato cloud and the legacy MPLS network co-exist behind a datacenter socket until the last circuit is switched off. No big-bang cutover, no service disruption, and every migrated site lands on the network design you actually want.

[MPLS](#)[Migration](#)[Co-existence](#)[Routing](#)

BUSINESS OBJECTIVE

Why this matters

Most organisations that buy Cato are not building a WAN from scratch — they are leaving one. MPLS contracts stagger across regions, some sites have circuits with years left to run, and the applications in the datacenter must stay reachable from every site on every day of the migration. The deck behind this page states the objectives plainly:

- **Co-existence during migration** — Cato and MPLS run side by side, for as long as needed.
- **From MPLS to an elegant future network design** — migrate to the target architecture, not a copy of the old one.
- **Seamless integration** — the datacenter bridges both worlds with no application changes.
- **Minimise service disruption** — each site cuts over on its own schedule, with dual last miles from day one.
- **New sites talk to current sites** — full any-to-any reachability between migrated and legacy sites, throughout.

Why WAN migrations stall

The big-bang trap

Flag-day cutovers concentrate all the risk in one change window. One misrouted prefix and every site loses the datacenter — so the migration gets postponed, again.

Contract overlap

MPLS terms rarely end on the same date. Sites need to leave as their circuits expire — which means the interim network must work for months, not days.

The in-between network

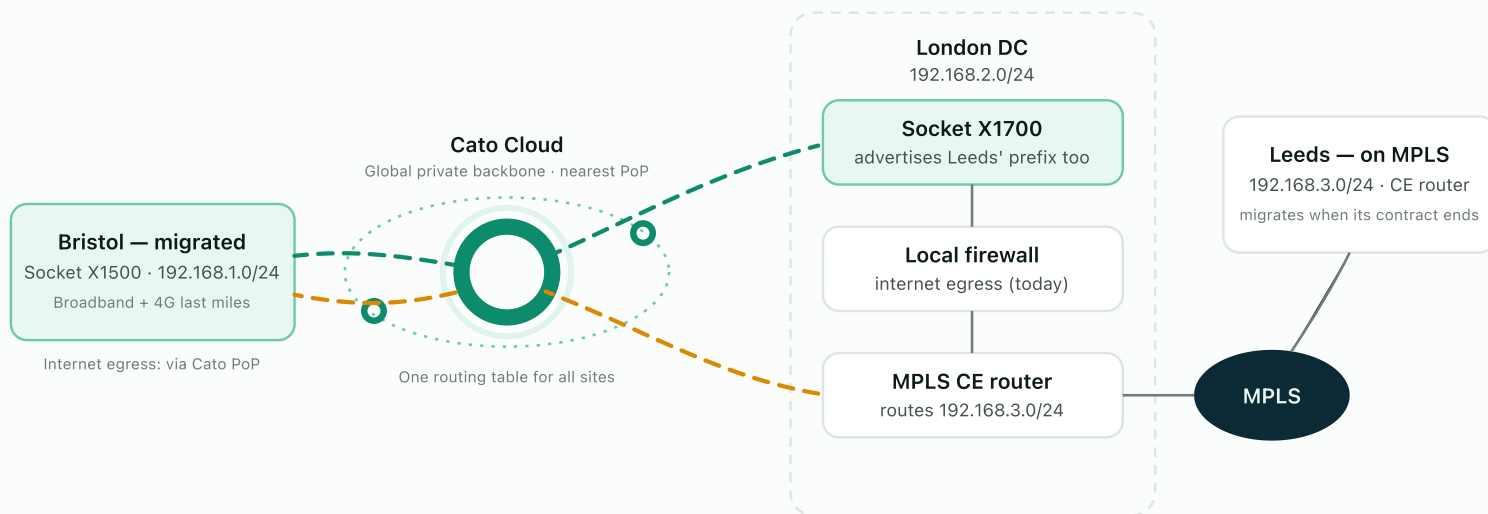
The hard part isn't either end state — it's the middle: who routes between the new WAN and the old one, and who can prove it's working at any point in time?

Discovery question

"When your MPLS renewal lands, which sites could you move tomorrow — and which are stuck because other sites still need to reach them? Who owns the routing design for the in-between state, and how long do you have to run it?"

Co-existence by design

The pattern is simple and repeatable. The datacenter gets a socket first — it becomes the bridge between both networks. Its socket advertises not only the local DC prefix but also the prefixes of every site still on MPLS, reached via the DC LAN and the existing MPLS CE router. Each branch then migrates on its own schedule: drop in a socket with broadband and 4G last miles, and it is instantly on-net to the DC *and* — transiting the DC — to every site still on MPLS. Nothing about the legacy network changes until its circuit is switched off.



--- Bristol ↔ London DC — on-net via Cato -.- Bristol ↔ Leeds — transits London DC during co-existence

— Legacy MPLS — unchanged until Leeds migrates

The London DC socket advertises Leeds' prefix into the Cato cloud, so migrated sites reach un-migrated sites via the DC LAN and the existing MPLS router — no changes to the legacy network, no big-bang.



Co-existence by design

Any prefix behind any socket is routable — including MPLS-site prefixes reached via a static route at the DC socket. New and current sites talk throughout.



Site-by-site cutover

Sockets are zero-touch: ship, plug in broadband and 4G, and the site is on-net. Each site moves as its MPLS term expires — no synchronised change windows.



An elegant end state

You migrate to the target design — full-mesh any-to-any over the global backbone with optimisation built in — not a re-implementation of hub-and-spoke MPLS.

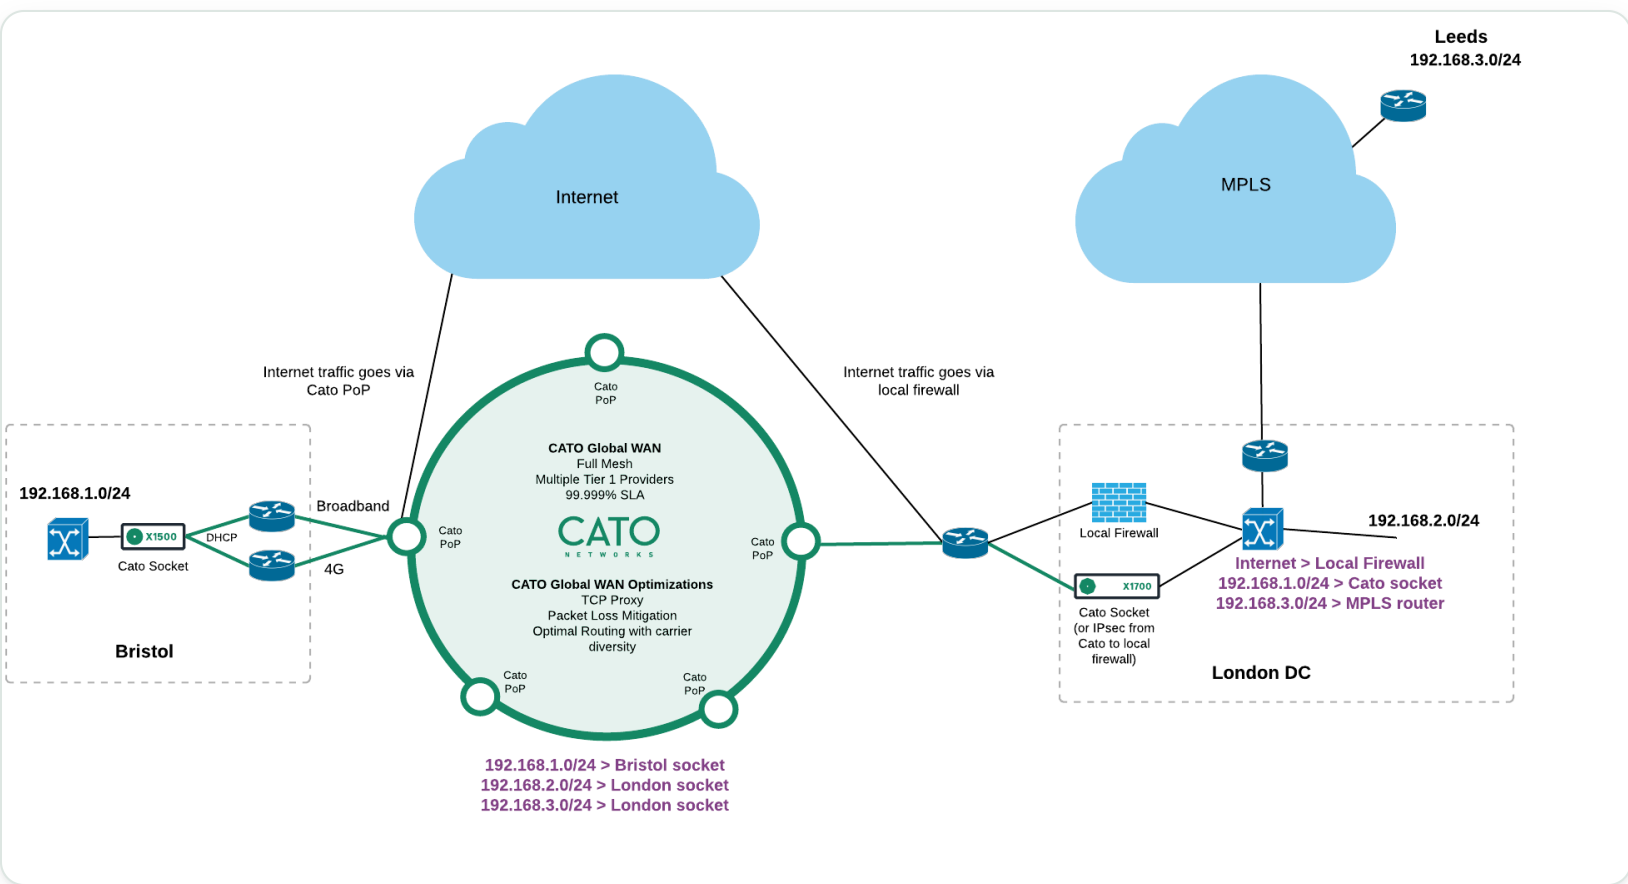


Minimal disruption

Dual last miles from day one, and internet egress can stay on the DC firewall until you're ready to move it — users notice nothing during the transition.

Who routes what during the transition

The original deck diagram below shows the full picture: Bristol already migrated (X1500, broadband + 4G, internet breakout at the Cato PoP), London DC bridging both networks (X1700 socket — or IPsec from Cato to the local firewall — plus the firewall and MPLS router), and Leeds still on MPLS.



Deck original: co-existence during migration — Bristol on Cato, Leeds on MPLS, London DC bridging both, with the routing tables that make it work.

The two routing tables in that diagram are the whole trick — worth walking line by line in the room:

DESTINATION	SITE	CATO ROUTING TABLE	LONDON DC LOCAL ROUTING
192.168.1.0/24	Bristol (migrated)	→ Bristol socket	→ Cato socket
192.168.2.0/24	London DC	→ London socket	local LAN
192.168.3.0/24	Leeds (still on MPLS)	→ London socket	→ MPLS CE router
Internet	—	Bristol: breakout at Cato PoP	→ local firewall

As each MPLS site migrates, its prefix simply moves from "→ London socket" to its own socket in the Cato routing table, and the corresponding static route is deleted at the DC. When the last site moves, the MPLS router and the contract behind it are decommissioned.

How to run the demo

Note — authored runbook

The “How to” slide in the original deck was blank, so this runbook was authored for the library from the co-existence topology above and the standard demo tenant. Review it against your tenant before first customer use.

1 Walk the co-existence topology

Open on the diagram above. Three sites, three states: Bristol already migrated, Leeds untouched on MPLS, London DC bridging both. Ask which of their sites map to each role — most prospects can name their “London DC” immediately.

2 Show how the DC socket advertises MPLS-site routes

Network → Sites

Open London DC and show its networks and static routes: the local DC range, plus the Leeds prefix (192.168.3.0/24) with a next hop of the MPLS CE router on the DC LAN. This is the entire integration — no BGP redesign, no changes on the MPLS side.

3 Show the routing table both networks share

Monitor → Topology

Show every site, its tunnels and its ranges in one view — migrated prefixes behind their own sockets, the Leeds prefix behind London DC. If you have Socket WebUI access, the socket's routing table tells the same story from the site's point of view.

4 Simulate the cutover of one site

Network → Sites

Create a new socket site to play the part of Leeds on cutover day (use a spare lab socket, e.g. the Bracknell Lab X1500, or just walk the site-creation flow): assign the 192.168.3.0/24 range to the new site, delete the static route at London DC, and the Cato routing table updates. That is the whole change window.

5 Prove new and current sites keep talking

[Monitor → Topology](#)

Generate traffic between a migrated site and the DC (file share, RDP, ping) and show the live tunnels and throughput. Then [Monitor → App Analytics](#) for the same flows attributed to sites and applications — the visibility they never had on MPLS.

6 Discuss internet egress during the transition

[Security → Internet Firewall](#)

Two valid interim patterns, per site: keep breaking out via the DC firewall (nothing changes for compliance or SIEM feeds), or break out directly at the Cato PoP with the full security stack — as Bristol does in the diagram. Show that this is a per-site policy choice, not an architectural fork.

7 Close on the end state

When the last site migrates: delete the static route, decommission the MPLS router, terminate the contract. What remains is the elegant target design — every site any-to-any on the global backbone, one policy, one console.

GO DEEPER

Resources

Documentation

- [Cato SASE Cloud architecture — platform page](#)
- [Cato Management Application \(CMA\)](#)
- [Cato knowledge base — sockets, sites & routing](#)

In this library

- [SD-WAN: Global Private Access](#)
- [Cloud Datacenter Integration](#)